
Deep Learning (MATLAB)

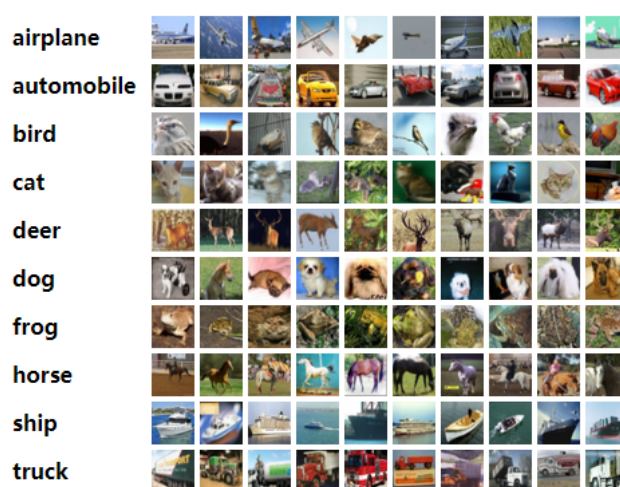
Contents

1	The CIFAR-10 dataset	2
2	Convolutional Neural Networks (CNNs / ConvNets)	2
2.1	Architecture Overview	2
2.2	Layers used to build ConvNets	4
2.2.1	Convolutional Layer	4
2.2.2	Pooling Layer	5
3	Deep Learning Softwares	7
4	Tasks	7
5	Codes	7
6	Results	13
7	Reference	13

1 The CIFAR-10 dataset

The CIFAR-10 dataset (<http://www.cs.toronto.edu/~kriz/cifar.html>) consists of 60000 32×32 colour images in 10 classes, with 6000 images per class. There are 50000 training images and 10000 test images.

The dataset is divided into five training batches and one test batch, each with 10000 images. The test batch contains exactly 1000 randomly-selected images from each class. The training batches contain the remaining images in random order, but some training batches may contain more images from one class than another. Between them, the training batches contain exactly 5000 images from each class. Here are the classes in the dataset, as well as 10 random images from each:



The classes are completely mutually exclusive. There is no overlap between automobiles and trucks. "Automobile" includes sedans, SUVs, things of that sort. "Truck" includes only big trucks. Neither includes pickup trucks.


2 Convolutional Neural Networks (CNNs / ConvNets)

2.1 Architecture Overview

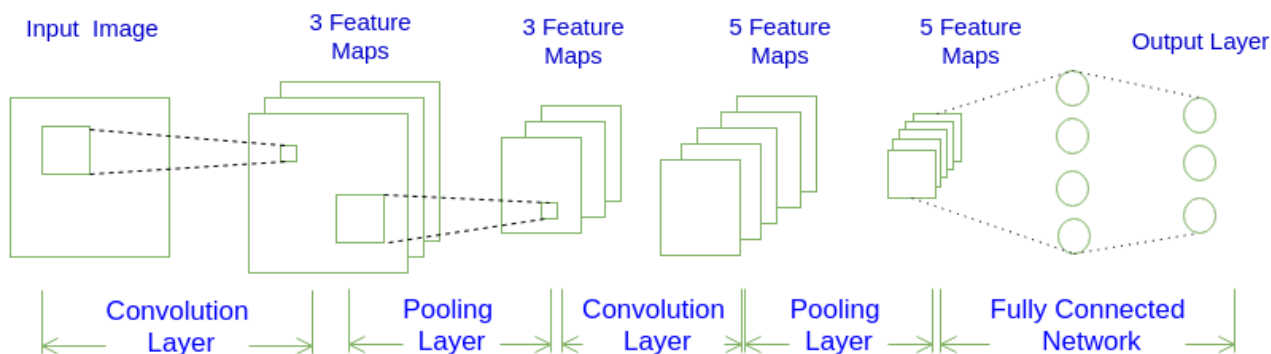
Regular Neural Nets don't scale well to full images. In CIFAR-10, images are only of size $32 \times 32 \times 3$ (32 wide, 32 high, 3 color channels), so a single fully-connected neuron in a first hidden layer of a regular Neural Network would have $32 * 32 * 3 = 3072$ weights. This amount still seems manageable, but clearly this fully-connected structure does not scale to larger images. For example, an image of more respectable size, e.g. $200 \times 200 \times 3$, would lead to neurons that have $200 * 200 * 3 = 120,000$ weights. Moreover, we would almost certainly want to have several such neurons, so the parameters

would add up quickly! Clearly, this full connectivity is wasteful and the huge number of parameters would quickly lead to overfitting.

Convolutional Neural Networks take advantage of the fact that the input consists of images and they constrain the architecture in a more sensible way. In particular, unlike a regular Neural Network, the layers of a ConvNet have neurons arranged in 3 dimensions: width, height, depth. (Note that the word depth here refers to the third dimension of an activation volume, not to the depth of a full Neural Network, which can refer to the total number of layers in a network.) For example, the input images in CIFAR-10 are an input volume of activations, and the volume has dimensions $32 \times 32 \times 3$ (width, height, depth respectively). As we will soon see, the neurons in a layer will only be connected to a small region of the layer before it, instead of all of the neurons in a fully-connected manner. Moreover, the final output layer would for CIFAR-10 have dimensions $1 \times 1 \times 10$, because by the end of the ConvNet architecture we will reduce the full image into a single vector of class scores, arranged along the depth dimension. Here is a visualization:



Left: A regular 3-layer Neural Network. Right: A ConvNet arranges its neurons in three dimensions (width, height, depth), as visualized in one of the layers. Every layer of a ConvNet transforms the 3D input volume to a 3D output volume of neuron activations. In this example, the red input layer holds the image, so its width and height would be the dimensions of the image, and the depth would be 3 (Red, Green, Blue channels).



2.2 Layers used to build ConvNets

a simple ConvNet is a sequence of layers, and every layer of a ConvNet transforms one volume of activations to another through a differentiable function. We use three main types of layers to build ConvNet architectures: **Convolutional Layer**, **Pooling Layer**, and **Fully-Connected Layer** (exactly as seen in regular Neural Networks). We will stack these layers to form a full ConvNet architecture.

Example Architecture: Overview. We will go into more details below, but a simple ConvNet for CIFAR-10 classification could have the architecture [INPUT - CONV - RELU - POOL - FC]. In more detail:

- INPUT $[32 \times 32 \times 3]$ will hold the raw pixel values of the image, in this case an image of width 32, height 32, and with three color channels R,G,B.
- CONV layer will compute the output of neurons that are connected to local regions in the input, each computing a dot product between their weights and a small region they are connected to in the input volume. This may result in volume such as $[32 \times 32 \times 12]$ if we decided to use 12 filters.
- RELU layer will apply an elementwise activation function, such as the $\max(0, x)$ thresholding at zero. This leaves the size of the volume unchanged ($[32 \times 32 \times 12]$).
- POOL layer will perform a downsampling operation along the spatial dimensions (width, height), resulting in volume such as $[16 \times 16 \times 12]$.
- FC (i.e. fully-connected) layer will compute the class scores, resulting in volume of size $[1 \times 1 \times 10]$, where each of the 10 numbers correspond to a class score, such as among the 10 categories of CIFAR-10. As with ordinary Neural Networks and as the name implies, each neuron in this layer will be connected to all the numbers in the previous volume.

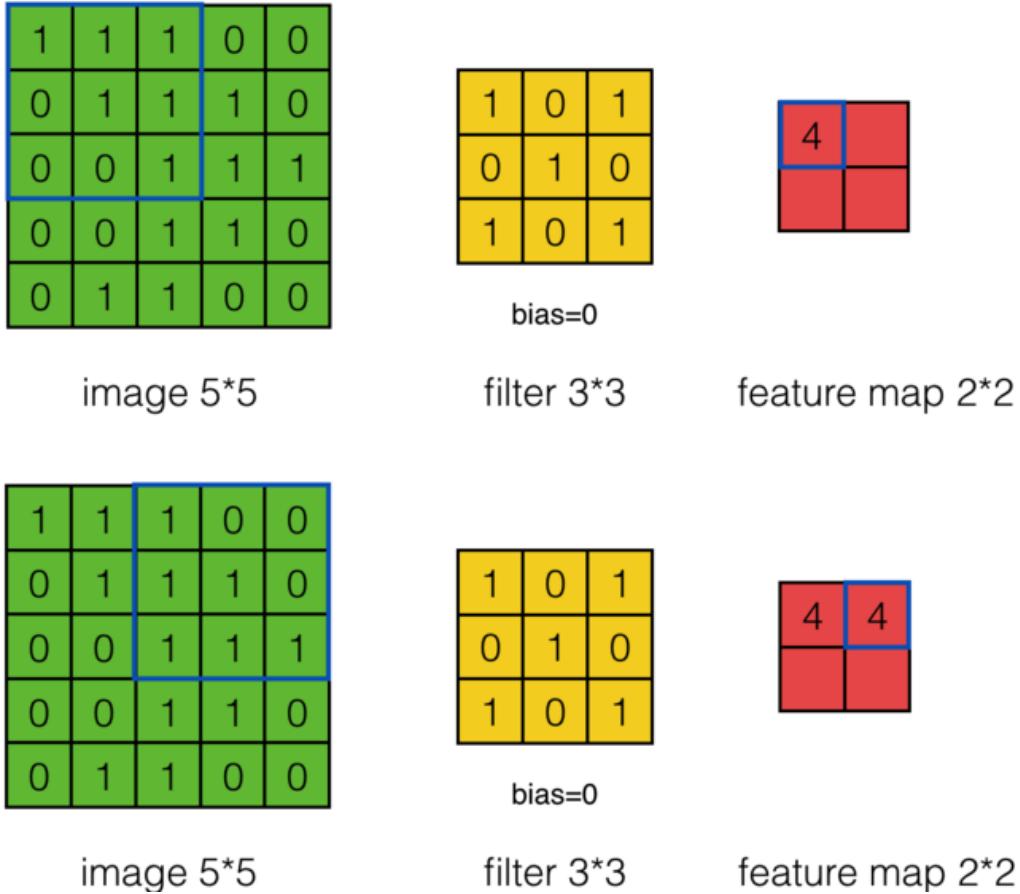
2.2.1 Convolutional Layer

To summarize, the Conv Layer:

- Accepts a volume of size $W_1 \times H_1 \times D_1$
- Requires four hyperparameters:
 - Number of filters K ,
 - their spatial extent F ,
 - the stride S ,
 - the amount of zero padding P .

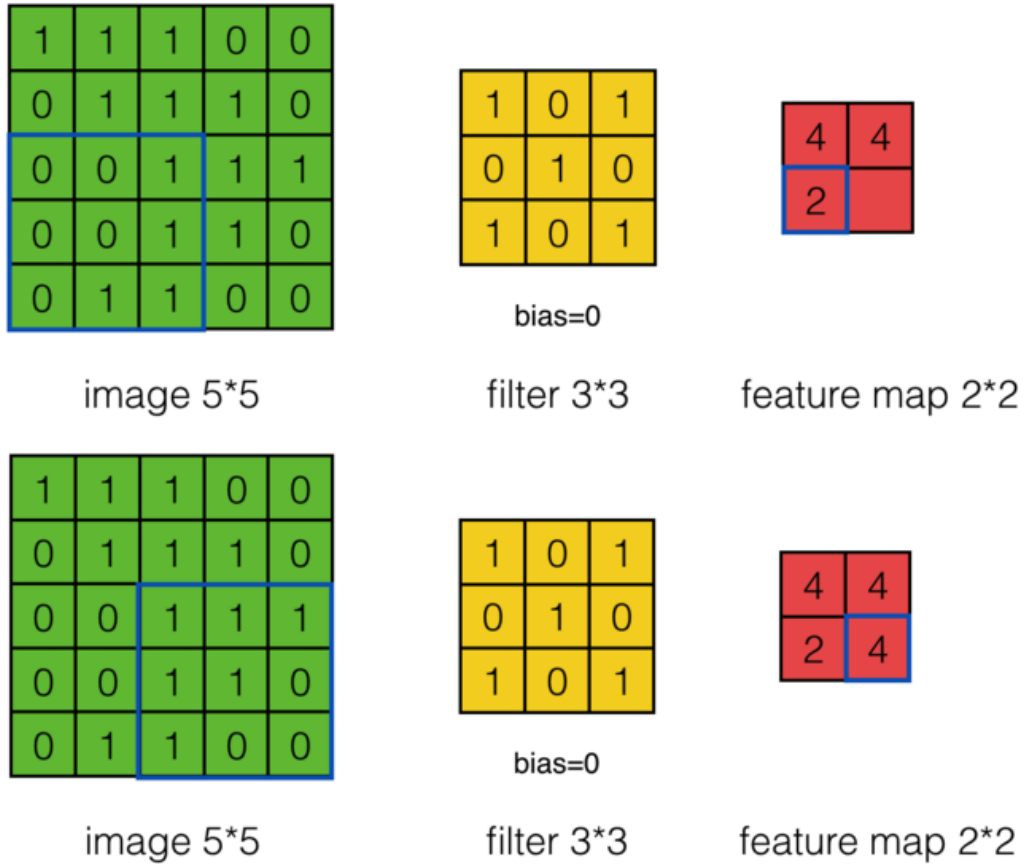
- Produces a volume of size $W_2 \times H_2 \times D_2$ where:
 - $W_2 = (W_1 - F + 2P)/S + 1$
 - $H_2 = (H_1 - F + 2P)/S + 1$ (i.e. width and height are computed equally by symmetry)
 - $D_2 = K$
- With parameter sharing, it introduces $F \cdot F \cdot D_1$ weights per filter, for a total of $(F \cdot F \cdot D_1) \cdot K$ weights and K biases.
- In the output volume, the d -th depth slice (of size $W_2 \times H_2$) is the result of performing a valid convolution of the d -th filter over the input volume with a stride of S , and then offset by d -th bias.

A common setting of the hyperparameters is $F = 3, S = 1, P = 1$. However, there are common conventions and rules of thumb that motivate these hyperparameters.



2.2.2 Pooling Layer

It is common to periodically insert a Pooling layer in-between successive Conv layers in a ConvNet architecture. Its function is to progressively reduce the spatial size of the representation to reduce the

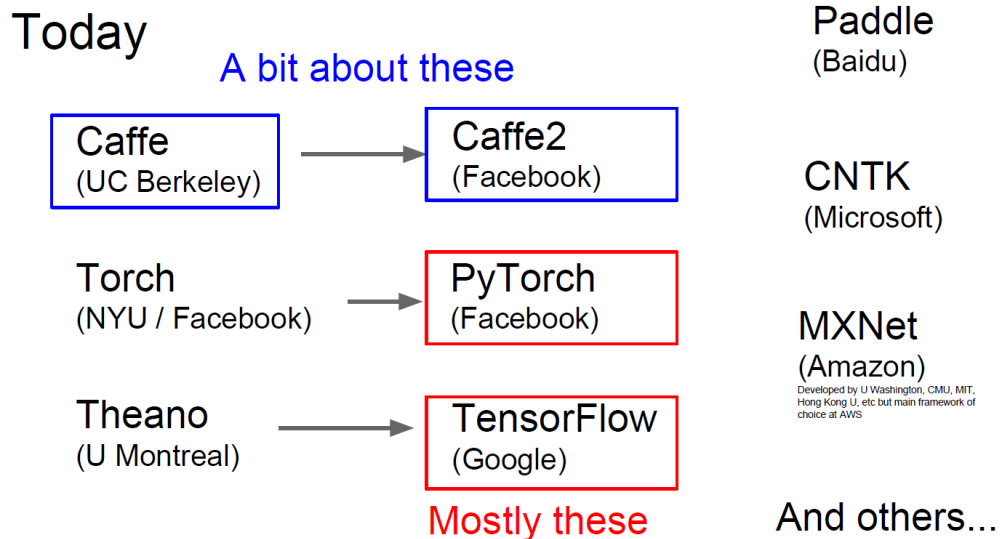


amount of parameters and computation in the network, and hence to also control overfitting. The Pooling Layer operates independently on every depth slice of the input and resizes it spatially, using the **MAX** operation. The most common form is a pooling layer with filters of size 2×2 applied with a stride of 2 downsamples every depth slice in the input by 2 along both width and height, discarding 75% of the activations. Every MAX operation would in this case be taking a max over 4 numbers (little 2×2 region in some depth slice). The depth dimension remains unchanged. More generally, the pooling layer:

- Accepts a volume of size $W_1 \times H_1 \times D_1$
- Requires two hyperparameters:
 - their spatial extent F ,
 - the stride S ,
- Produces a volume of size $W_2 \times H_2 \times D_2$ where:
 - $W_2 = (W_1 - F)/S + 1$
 - $H_2 = (H_1 - F)/S + 1$
 - $D_2 = D_1$

- Introduces zero parameters since it computes a fixed function of the input
- For Pooling layers, it is not common to pad the input using zero-padding.

3 Deep Learning Softwares



4 Tasks

1. Given the data set in the first section, please implement a convolutional neural network to calculate the accuracy rate. The major steps involved are as follows:
 - (a) Reading the input image.
 - (b) Preparing filters.
 - (c) Conv layer: Convoluting each filter with the input image.
 - (d) ReLU layer: Applying ReLU activation function on the feature maps (output of conv layer).
 - (e) Max Pooling layer: Applying the pooling operation on the output of ReLU layer.
 - (f) Stacking conv, ReLU, and max pooling layers
2. You can refer to the codes in `cs231n`. Don't use Keras, TensorFlow, PyTorch, Theano, Caffe, and other deep learning softwares.

5 Codes

```

1  clear;
2  [trainingImages,trainingLabels,testImages,testLabels] = helperCIFAR10Data.load('');
3  numImageCategories = 10;
4  % Create the image input layer for 32x32x3 CIFAR-10 images.
5  [height,width,numChannels, ~] = size(trainingImages);
6
7  imageSize = [height width numChannels];
8  inputLayer = imageInputLayer(imageSize)
9  % Convolutional layer parameters
10 filterSize = [5 5];
11 numFilters = 32;
12
13 middleLayers = [
14
15 % The first convolutional layer has a bank of 32 5x5x3 filters. A
16 % symmetric padding of 2 pixels is added to ensure that image borders
17 % are included in the processing. This is important to avoid
18 % information at the borders being washed away too early in the
19 % network.
20 convolution2dLayer(filterSize,numFilters,'Padding',2)
21
22 % Note that the third dimension of the filter can be omitted because it
23 % is automatically deduced based on the connectivity of the network. In
24 % this case because this layer follows the image layer, the third
25 % dimension must be 3 to match the number of channels in the input
26 % image.
27
28 % Next add the ReLU layer:
29 reluLayer()
30
31 % Follow it with a max pooling layer that has a 3x3 spatial pooling area
32 % and a stride of 2 pixels. This down-samples the data dimensions from
33 % 32x32 to 15x15.
34 maxPooling2dLayer(3,'Stride',2)
35
36 % Repeat the 3 core layers to complete the middle of the network.
37 convolution2dLayer(filterSize,numFilters,'Padding',2)
38 reluLayer()
39 maxPooling2dLayer(3, 'Stride',2)
40

```



```

41 convolution2dLayer(filterSize,2 * numFilters,'Padding',2)
42 reluLayer()
43 maxPooling2dLayer(3,'Stride',2)
44
45 ]% Convolutional layer parameters
46 filterSize = [5 5];
47 numFilters = 32;
48
49 middleLayers = [
50
51 % The first convolutional layer has a bank of 32 5x5x3 filters. A
52 % symmetric padding of 2 pixels is added to ensure that image borders
53 % are included in the processing. This is important to avoid
54 % information at the borders being washed away too early in the
55 % network.
56 convolution2dLayer(filterSize,numFilters,'Padding',2)
57
58 % Note that the third dimension of the filter can be omitted because it
59 % is automatically deduced based on the connectivity of the network. In
60 % this case because this layer follows the image layer, the third
61 % dimension must be 3 to match the number of channels in the input
62 % image.
63
64 % Next add the ReLU layer:
65 reluLayer()
66
67 % Follow it with a max pooling layer that has a 3x3 spatial pooling area
68 % and a stride of 2 pixels. This down-samples the data dimensions from
69 % 32x32 to 15x15.
70 maxPooling2dLayer(3,'Stride',2)
71
72 % Repeat the 3 core layers to complete the middle of the network.
73 convolution2dLayer(filterSize,numFilters,'Padding',2)
74 reluLayer()
75 maxPooling2dLayer(3,'Stride',2)
76
77 convolution2dLayer(filterSize,2 * numFilters,'Padding',2)
78 reluLayer()
79 maxPooling2dLayer(3,'Stride',2)
80
81 ]

```

```

82
83
84 finalLayers = [
85
86 % Add a fully connected layer with 64 output neurons. The output size of
87 % this layer will be an array with a length of 64.
88 fullyConnectedLayer(64)
89
90 % Add an ReLU non-linearity.
91 reluLayer
92
93 % Add the last fully connected layer. At this point, the network must
94 % produce 10 signals that can be used to measure whether the input image
95 % belongs to one category or another. This measurement is made using the
96 % subsequent loss layers.
97 fullyConnectedLayer(numImageCategories)
98
99 % Add the softmax loss layer and classification layer. The final layers use
100 % the output of the fully connected layer to compute the categorical
101 % probability distribution over the image classes. During the training
102 % process, all the network weights are tuned to minimize the loss over this
103 % categorical distribution.
104 softmaxLayer
105 classificationLayer
106 ]
107
108 layers = [
109     inputLayer
110     middleLayers
111     finalLayers
112 ]
113
114 layers(2).Weights = 0.0001 * randn([filterSize numChannels numFilters]);
115
116 % Set the network training options
117 opts = trainingOptions('sgdm', ...
118     'Momentum', 0.9, ...
119     'InitialLearnRate', 0.001, ...
120     'LearnRateSchedule', 'piecewise', ...
121     'LearnRateDropFactor', 0.1, ...
122     'LearnRateDropPeriod', 8, ...

```

```

123     'L2Regularization', 0.004, ...
124     'MaxEpochs', 40, ...
125     'MiniBatchSize', 128, ...
126     'Verbose', true);
127
128 % A trained network is loaded from disk to save time when running the
129 % example. Set this flag to true to train the network.
130 doTraining = false;
131
132 if doTraining
133     % Train a network.
134     cifar10Net = trainNetwork(trainingImages, trainingLabels, layers, opts);
135 else
136     % Load pre-trained detector for the example.
137     load('rcnnStopSigns.mat', 'cifar10Net')
138 end
139
140 % Extract the first convolutional layer weights
141 w = cifar10Net.Layers(2).Weights;
142
143 % rescale the weights to the range [0, 1] for better visualization
144 w = rescale(w);
145
146 figure
147 montage(w)
148
149 % Run the network on the test set.
150 YTest = classify(cifar10Net, testImages);
151
152 % Calculate the accuracy.
153 accuracy = sum(YTest == testLabels)/numel(testLabels)
154
155 % Load the ground truth data
156 data = load('stopSignsAndCars.mat', 'stopSignsAndCars');
157 stopSignsAndCars = data.stopSignsAndCars;
158
159 % Update the path to the image files to match the local file system
160 visiondata = fullfile(toolboxdir('vision'),'visiondata');
161 stopSignsAndCars.imageFilename = fullfile(visiondata, stopSignsAndCars.imageFilename);
162
163 % Display a summary of the ground truth data

```

```

164 summary(stopSignsAndCars)
165
166 % Only keep the image file names and the stop sign ROI labels
167 stopSigns = stopSignsAndCars(:, {'imageFilename','stopSign'});
168
169 % Display one training image and the ground truth bounding boxes
170 I = imread(stopSigns.imageFilename{1});
171 I = insertObjectAnnotation(I,'Rectangle',stopSigns.stopSign{1},'stop_sign','LineWidth',8);
172
173 figure
174 imshow(I)
175 % A trained detector is loaded from disk to save time when running the
176 % example. Set this flag to true to train the detector.
177 doTraining = false;
178
179 if doTraining
180
181     % Set training options
182     options = trainingOptions('sgdm', ...
183         'MiniBatchSize', 128, ...
184         'InitialLearnRate', 1e-3, ...
185         'LearnRateSchedule', 'piecewise', ...
186         'LearnRateDropFactor', 0.1, ...
187         'LearnRateDropPeriod', 100, ...
188         'MaxEpochs', 100, ...
189         'Verbose', true);
190
191     % Train an R-CNN object detector. This will take several minutes.
192     rcnn = trainRCNNObjectDetector(stopSigns, cifar10Net, options, ...
193         'NegativeOverlapRange', [0 0.3], 'PositiveOverlapRange',[0.5 1])
194 else
195     % Load pre-trained network for the example.
196     load('rcnnStopSigns.mat','rcnn')
197 end
198 % Read test image
199 testImage = imread('stopSignTest.jpg');
200
201 % Detect stop signs
202 [bboxes,score,label] = detect(rcnn,testImage,'MiniBatchSize',128)
203

```

```

204
205 % Display the detection results
206 [score, idx] = max(score);
207
208 bbox = bboxes(idx, :);
209 annotation = sprintf('s: (Confidence = %f)', label(idx), score);
210
211 outputImage = insertObjectAnnotation(testImage, 'rectangle', bbox, annotation);
212
213 figure
214 imshow(outputImage)
215 % The trained network is stored within the R-CNN detector
216 rcnn.Network
217
218
219 featureMap = activations(rcnn.Network, testImage, 14);
220
221 % The softmax activations are stored in a 3-D array.
222 size(featureMap)
223 rcnn.ClassNames
224
225 stopSignMap = featureMap(:, :, 1);
226 % Resize stopSignMap for visualization
227 [height, width, ~] = size(testImage);
228 stopSignMap = imresize(stopSignMap, [height, width]);
229
230 % Visualize the feature map superimposed on the test image.
231 featureMapOnImage = imfuse(testImage, stopSignMap);
232
233 figure
234 imshow(featureMapOnImage)

```

6 Results

7 Reference

- <https://www.zybuluo.com/hanbingtao/note/485480>
- <http://cs231n.github.io/convolutional-networks/#layers>

```
inputLayer =
```

[ImageInputLayer](#) - 属性:

```
        Name: ''
    InputSize: [32 32 3]
```

超参数

```
    DataAugmentation: 'none'
    Normalization: 'zerocenter'
    NormalizationDimension: 'auto'
    Mean: []
```

Figure 1: input layer

```
middleLayers =
```

具有以下层的 9x1 [Layer](#) 数组:

1	''	卷积	32 5x5 卷积: 步幅 [1 1], 填充 [2 2 2 2]
2	''	ReLU	ReLU
3	''	最大池化	3x3 最大池化: 步幅 [2 2], 填充 [0 0 0 0]
4	''	卷积	32 5x5 卷积: 步幅 [1 1], 填充 [2 2 2 2]
5	''	ReLU	ReLU
6	''	最大池化	3x3 最大池化: 步幅 [2 2], 填充 [0 0 0 0]
7	''	卷积	64 5x5 卷积: 步幅 [1 1], 填充 [2 2 2 2]
8	''	ReLU	ReLU
9	''	最大池化	3x3 最大池化: 步幅 [2 2], 填充 [0 0 0 0]

Figure 2: middle layer

```
layers =
```

具有以下层的 15x1 [Layer](#) 数组:

1	''	图像输入	32x32x3 图像: 'zerocenter' 归一化
2	''	卷积	32 5x5 卷积: 步幅 [1 1], 填充 [2 2 2 2]
3	''	ReLU	ReLU
4	''	最大池化	3x3 最大池化: 步幅 [2 2], 填充 [0 0 0 0]
5	''	卷积	32 5x5 卷积: 步幅 [1 1], 填充 [2 2 2 2]
6	''	ReLU	ReLU
7	''	最大池化	3x3 最大池化: 步幅 [2 2], 填充 [0 0 0 0]
8	''	卷积	64 5x5 卷积: 步幅 [1 1], 填充 [2 2 2 2]
9	''	ReLU	ReLU
10	''	最大池化	3x3 最大池化: 步幅 [2 2], 填充 [0 0 0 0]
11	''	全连接	64 全连接层
12	''	ReLU	ReLU
13	''	全连接	10 全连接层
14	''	Softmax	softmax
15	''	分类输出	crossentropyex

Figure 3: layers

```
accuracy =
```

0.7456

Figure 4: accuracy