

Deep Learning

School of Data and Computer Science
Sun Yat-sen University



- 1 Datasets
- 2 Convolutional Neural Networks (CNNs / ConvNets)
- 3 Deep Learning Softwares
- 4 Tasks



1 Datasets

2 Convolutional Neural Networks (CNNs / ConvNets)

3 Deep Learning Softwares

4 Tasks



The CIFAR-10 dataset

- The CIFAR-10 dataset
(<http://www.cs.toronto.edu/~kriz/cifar.html>)
consists of 60000 32×32 colour images in 10 classes, with 6000 images per class. There are 50000 training images and 10000 test images.
- The dataset is divided into five training batches and one test batch, each with 10000 images.
- The test batch contains exactly 1000 randomly-selected images from each class.



The CIFAR-10 dataset (cont'd)

- The training batches contain the remaining images in random order, but some training batches may contain more images from one class than another.
- Between them, the training batches contain exactly 5000 images from each class.
- The classes are completely mutually exclusive. There is no overlap between automobiles and trucks.
 - ◊ "Automobile" includes sedans, SUVs, things of that sort.
 - ◊ "Truck" includes only big trucks.
 - ◊ Neither includes pickup trucks.



The CIFAR-10 dataset (cont'd)

Here are the classes in the dataset, as well as 10 random images from each:

airplane




automobile



bird



cat




deer



dog




frog



horse



ship



truck



- 1 Datasets
- 2 Convolutional Neural Networks (CNNs / ConvNets)
- 3 Deep Learning Softwares
- 4 Tasks



Architecture Overview

- Regular Neural Nets do not scale well to full images.
- Convolutional Neural Networks take advantage of the fact that the input consists of images and they constrain the architecture in a more sensible way.
- The layers of a ConvNet have neurons arranged in 3 dimensions: width, height, depth.
- For example, the input images in CIFAR-10 are an input volume of activations, and the volume has dimensions $32 \times 32 \times 3$ (width, height, depth respectively).



Architecture Overview (cont'd)

The final output layer would for CIFAR-10 have dimensions $1 \times 1 \times 10$. Here is a visualization:




Figure 1: *

Left: A regular 3-layer Neural Network.

Right: A ConvNet arranges its neurons in three dimensions (width, height, depth)



Layers used to build ConvNets

- A simple ConvNet is a sequence of layers, and every layer of a ConvNet transforms one volume of activations to another through a differentiable function.
- We use three main types of layers to build ConvNet architectures:
 - ◊ **Convolutional Layer**
 - ◊ **Pooling Layer**
 - ◊ **Fully-Connected Layer**
- We will stack these layers to form a full ConvNet architecture.



Layers used to build ConvNets (cont'd)




Figure 2: *Example Architecture*: Overview.



ConvNet for CIFAR-10 classification

A simple ConvNet for CIFAR-10 classification could have the architecture [**INPUT** - **CONV** - **RELU** - **POOL** - **FC**].

- **INPUT** [$32 \times 32 \times 3$] will hold the raw pixel values of the image, in this case an image of width 32, height 32, and with three color channels R,G,B.
- **CONV** layer will compute the output of neurons that are connected to local regions in the input, each computing a dot product between their weights and a small region they are connected to in the input volume. This may result in volume such as [$32 \times 32 \times 12$] if we decided to use 12 filters.
- **RELU** layer will apply an elementwise activation function, such as the $\max(0, x)$ thresholding at zero. This leaves the size of the volume unchanged ($[32 \times 32 \times 12]$).



ConvNet for CIFAR-10 classification (cont'd)

- **POOL** layer will perform a downsampling operation along the spatial dimensions (width, height), resulting in volume such as $[16 \times 16 \times 12]$.
- **FC** (i.e. fully-connected) layer will compute the class scores, resulting in volume of size $[1 \times 1 \times 10]$, where each of the 10 numbers correspond to a class score, such as among the 10 categories of CIFAR-10. As with ordinary Neural Networks and as the name implies, each neuron in this layer will be connected to all the numbers in the previous volume.



Layers used to build ConvNets

Convolutional Layer

To summarize, the Conv Layer:

- Accepts a volume of size $W_1 \times H_1 \times D_1$
- Requires four hyperparameters:
 - Number of filters K ,
 - their spatial extent F ,
 - the stride S ,
 - the amount of zero padding P .
- Produces a volume of size $W_2 \times H_2 \times D_2$ where:
 - $W_2 = (W_1 - F + 2P)/S + 1$
 - $H_2 = (H_1 - F + 2P)/S + 1$ (i.e. width and height are computed equally by symmetry)
 - $D_2 = K$



Layers used to build ConvNets

Convolutional Layer (cont'd)

- With parameter sharing, it introduces $F \cdot F \cdot D_1$ weights per filter, for a total of $(F \cdot F \cdot D_1) \cdot K$ weights and K biases.
- In the output volume, the d -th depth slice (of size $W_2 \times H_2$) is the result of performing a valid convolution of the d -th filter over the input volume with a stride of S , and then offset by d -th bias.



Layers used to build ConvNets

Convolutional Layer (cont'd)

A common setting of the hyperparameters is $F = 3, S = 1, P = 1$.

1	1	1	0	0
0	1	1	1	0
0	0	1	1	1
0	0	1	1	0
0	1	1	0	0

image 5*5

1	0	1
0	1	0
1	0	1

bias=0

filter 3*3

4	

feature map 2*2

1	1	1	0	0
0	1	1	1	0
0	0	1	1	1
0	0	1	1	0
0	1	1	0	0

image 5*5

1	0	1
0	1	0
1	0	1

bias=0

filter 3*3

4	4

feature map 2*2



Layers used to build ConvNets

Convolutional Layer (cont'd)

1	1	1	0	0
0	1	1	1	0
0	0	1	1	1
0	0	1	1	0
0	1	1	0	0

image 5*5

1	0	1
0	1	0
1	0	1

bias=0

filter 3*3

4	4
2	

feature map 2*2

1	1	1	0	0
0	1	1	1	0
0	0	1	1	1
0	0	1	1	0
0	1	1	0	0

image 5*5

1	0	1
0	1	0
1	0	1

bias=0

filter 3*3

4	4
2	4

feature map 2*2

Layers used to build ConvNets

Pooling Layer

- It is common to periodically insert a Pooling layer in-between successive Conv layers in a ConvNet architecture.
- Its function is to progressively reduce the spatial size of the representation to reduce the amount of parameters and computation in the network, and hence to also control overfitting.
- The Pooling Layer operates independently on every depth slice of the input and resizes it spatially, using the **MAX** operation.



Layers used to build ConvNets

Pooling Layer (cont'd)

- The most common form is a pooling layer with filters of size 2×2 applied with a stride of 2 downsamples every depth slice in the input by 2 along both width and height, discarding 75% of the activations.
- Every MAX operation would in this case be taking a max over 4 numbers (little 2×2 region in some depth slice).



Layers used to build ConvNets

Pooling Layer (cont'd)

The depth dimension remains unchanged. More generally, the pooling layer:

- Accepts a volume of size $W_1 \times H_1 \times D_1$
- Requires two hyperparameters:
 - their spatial extent F ,
 - the stride S ,
- Produces a volume of size $W_2 \times H_2 \times D_2$ where:
 - $W_2 = (W_1 - F)/S + 1$
 - $H_2 = (H_1 - F)/S + 1$
 - $D_2 = D_1$
- Introduces zero parameters since it computes a fixed function of the input
- For Pooling layers, it is not common to pad the input using zero-padding.




- 1 Datasets
- 2 Convolutional Neural Networks (CNNs / ConvNets)
- 3 Deep Learning Softwares
- 4 Tasks



Today

A bit about these



Mostly these

Paddle
(Baidu)

CNTK
(Microsoft)

MXNet
(Amazon)
Developed by U Washington, CMU, MIT,
Hong Kong U, etc but main framework of
choice at AWS

And others...



- Use **convolutional neural networks** (CNNs) and **long short-term memory** (LSTM) networks to perform classification and regression on image, time-series, and text data.
- Build advanced network architectures such as **generative adversarial networks** (GANs).
- Apps and plots can visualize **activations**, edit and analyze **network architectures**, and monitor **training progress**.
- Import models from **TensorFlow-Keras** and **Caffe**.
- Supports **transfer learning**.



MATLAB Deep Learning Toolbox

Create Simple Image Classification Network

1. Load Data

```
digitDatasetPath = fullfile(matlabroot,'toolbox','nnet','nndemos', ...  
    'nndatasets','DigitDataset');
```

```
imds = imageDatastore(digitDatasetPath, ...  
    'IncludeSubfolders',true, ...  
    'LabelSource','foldernames');
```

```
numTrainFiles = 750;    750 images
```

```
[imdsTrain,imdsValidation] = splitEachLabel(imds,numTrainFiles,'randomize'); training set and validation set
```

2. Define Network Architecture

```
inputSize = [28 28 1];    Each image is 28-by-28 pixels  
numClasses = 10;
```

```
layers = [  
    imageInputLayer(inputSize)  
    convolution2dLayer(5,20)  
    batchNormalizationLayer  
    reluLayer  
    fullyConnectedLayer(numClasses)  
    softmaxLayer  
    classificationLayer];
```

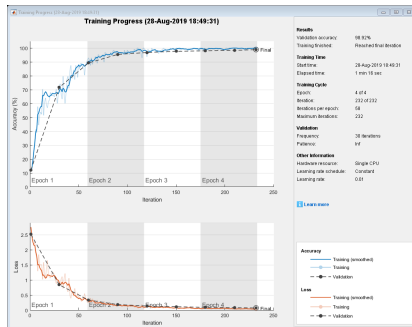


MATLAB Deep Learning Toolbox

Create Simple Image Classification Network (cont'd)

3. Train Network

```
options = trainingOptions('sgdm', ...  
    'MaxEpochs',4, ...  
    'ValidationData',imdsValidation, ...  
    'ValidationFrequency',30, ...  
    'Verbose',false, ...  
    'Plots','training-progress');  
  
net = trainNetwork(imdsTrain, layers, options);
```



4. Test Network

```
YPred = classify(net,imdsValidation);  
YValidation = imdsValidation.Labels;  
accuracy = mean(YPred == YValidation)
```

<https://ww2.mathworks.cn/help/deeplearning/gs/create-simple-deep-learning-classification-network.html>



- 1 Datasets
- 2 Convolutional Neural Networks (CNNs / ConvNets)
- 3 Deep Learning Softwares
- 4 Tasks



Given the dataset in the first section, please implement a convolutional neural network to calculate the accuracy rate with MATLAB. The major steps involved are as follows:

- ➊ Reading the input image.
- ➋ Preparing filters.
- ➌ Conv layer: Convolving each filter with the input image.
- ➍ ReLU layer: Applying ReLU activation function on the feature maps (output of conv layer).
- ➎ Max Pooling layer: Applying the pooling operation on the output of ReLU layer.
- ➏ Stacking conv, ReLU, and max pooling layers



- CIFAR-10 dataset,
<http://www.cs.toronto.edu/~kriz/cifar.html>
- <http://neuralnetworksanddeeplearning.com/chap6.html>
- CS231n Convolutional Neural Networks for Visual Recognition,
<https://cs231n.github.io/convolutional-networks/>
- MATLAB Deep Learning Toolbox,
<https://ww2.mathworks.cn/help/deeplearning/gs/create-simple-deep-learning-classification-network.html>



The End

